



## Growing-cache generation latency

Batch 1



Batch 8



2.836B parameters · BF16 · RTX 5090 · seeded random weights (no quality claim)

Cached last-token logits and KV.

2048-token prefix + 128 predetermined GPU-token decode calls. Sampling excluded; not serving latency. Bars: process 95% CI where n=3; otherwise trajectory round range, not CI.

After 10 attempts: retained A0001; frozen baseline A0000