



Accepted cumulative implementations — latency


Prefill, batch 1, length/context 2048




Decode, batch 1, length/context 2048



Prefill, batch 8, length/context 2048



Decode, batch 8, length/context 2048



2.836B parameters · BF16 · RTX 5090 · seeded random weights (no quality claim)

Cached last-token logits and KV; frozen baseline A0000.

Three independent balanced process blocks; 95% Student-t intervals (df=2). Baseline source is frozen A0000. Implementation history, not an additive decomposition.