



Fresh upstream / PR-head paired measurements

PR offload vs upstream main GPU



PR offload vs PR folded GPU



Peak allocated memory: main GPU minus PR offload

